



# STATISTICS FOR DATA SCIENCE

## Data Cleaning

---

**D. Uma**

Department of Computer Science and Engineering

# STATISTICS FOR DATA SCIENCE

---

## Data Cleaning

**D. Uma**

Department of Computer Science and Engineering

Suppose you have a dataset/database sitting in front of you, and I ask

**“Is it a good quality dataset/database?”**

This is **about the Data** themselves, not the system in use to access it.



### Data in the Real World is Dirty:


Lots of potentially incorrect data, e.g., instrument faulty, human or computer error, transmission error.

**Incomplete:** lacking attribute values, lacking certain attributes of interest, or containing only aggregate data.

e.g., *Occupation*=" " (missing data)

**Noisy:** containing noise, errors, or outliers

e.g., *Salary*="−10" (an error)



**Inconsistent:** containing discrepancies in codes or names, e.g.,

*Age="42", Birthday="03/07/2010"*

Was rating "1, 2, 3", now rating "A, B, C"

discrepancy between duplicate records

**Intentional** (e.g., *disguised missing data*)

Jan. 1 as everyone's birthday?





**Improved data quality** leads to **better decision making** across an organization.

The **more high-quality data** you have, the **more confidence** you can have in **your decisions**.

Good data **decreases risk** and can result in **consistent improvements** in **results**.

# STATISTICS FOR DATA SCIENCE

## Data Cleaning

**Data cleaning or cleansing** is the **process of detecting and correcting** (or removing) corrupt or inaccurate records from a record set, table, or database.

It also refers to **identifying incomplete, incorrect, inaccurate or irrelevant parts** of the data and then replacing, modifying, or deleting the dirty or coarse data.





- Makes the **data fit** for **purpose/plausible**
- **Reduces** the **negative impact of errors**
- **Improves** the **data quality**
- **Improves** the **quality of the outputs**

### PROCESS OF CLEANING DATA

- Detect
- Resolve
- Treat

- **Identify erroneous or suspicious data**
  - Graph or sort data - look at outliers
  - I have a **student who throws ten dice** and records the number of sixes. They recorded:  
**(2, 0, 3, 12, 2, 0, 1, 1, 3, 1, 4).**
  - What is wrong?
  - What do you think is the cause of it?

- Consider the data points
  - 3, 4, 7, 4, 8, 3, 9, 5, 7, 6, **92**
  - “**92**” is **suspicious** - an **outlier**
- Outliers:
  - are potentially legitimate (correct)
  - can be data or model glitches
  - can be a data miners dream, for example, a highly profitable customer
- **Outlier** - “**departure from the expected**”

### RESOLVE

- Deciding if **erroneous or suspicious data** should be **corrected** or amended
- Deciding on the action to “**treat**” the data

### WHAT TO LOOK FOR ?

- **Non-response**
  - an item non response
  - Eg. missing data
- **Erroneous data**
  - Can negatively affect data and resulting quality
- **Suspicious data**

- **Missing Data**
- **Irregular Data (Outliers)**
- **Unnecessary Data** — Repetitive Data, Duplicates and more
- **Inconsistent Data** — Capitalization, Addresses and more

# STATISTICS FOR DATA SCIENCE


## Handling Missing Data

### Technique : Missing Data Heat map

The chart demonstrates the missing data patterns of the **first 30 features**.

The horizontal axis shows the feature name; the vertical axis shows the number of observations/rows.

The **yellow color** represents the missing data while the blue color otherwise.





### Technique : Missing Data Percentage List

When there are **many features** in the dataset, we can make a list of **missing data %** for each feature.

This produces a list showing the percentage of missing values for each of the features.

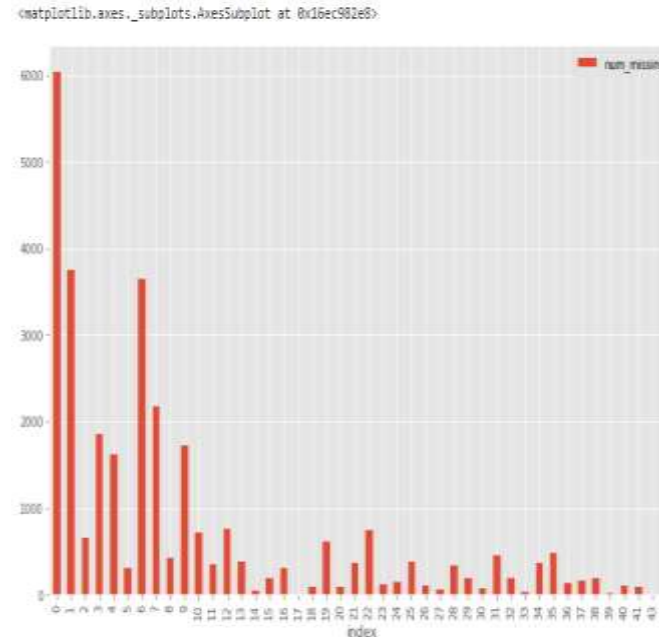
```
id - 0.0%
timestamp - 0.0%
full_sq - 0.0%
life_sq - 21.0%
floor - 1.0%
max_floor - 31.0%
material - 31.0%
build_year - 45.0%
num_room - 31.0%
kitch_sq - 31.0%
state - 44.0%
product_type - 0.0%
sub_area - 0.0%
area_m - 0.0%
raion_popul - 0.0%
green_zone_part - 0.0%
indust_part - 0.0%
children_preschool - 0.0%
preschool_quota - 22.0%
preschool_education_centers_raion - 0.0%
children_school - 0.0%
school_quota - 22.0%
school_education_centers_raion - 0.0%
school_education_centers_top_20_raion - 0.0%
hospital_beds_raion - 47.0%
healthcare_centers_raion - 0.0%
university_top_20_raion - 0.0%
sport_objects_raion - 0.0%
additional_education_raion - 0.0%
culture_objects_top_25 - 0.0%
```

# STATISTICS FOR DATA SCIENCE

## Handling Missing Data

### Technique : Missing Data Histogram

Missing data histogram is also a technique for when we have many features.



### Solution : Drop the Observation

In statistics, this method is called **the *listwise deletion technique***.

In this solution, we drop the entire observation as long as it contains a missing value.

Only if we are sure that the **missing data** is **not informative**, we perform this. Otherwise, we should consider other solutions.

### Solution : Drop the Feature

Similar to previous one, we only do this when we are confident that this feature doesn't provide useful information.

For example, from the missing data % list, we notice that `hospital_beds_raion` has a high missing value percentage of 47%. We may drop the entire feature.

Index	Age	Sex	Income
1	NA	M	NA
2	39	NA	75000
3	NA	NA	NA
4	28	F	50000
...	...	...	...
10000	18	F	NA

# STATISTICS FOR DATA SCIENCE

## Handling Missing Data

### Solution : Impute the Missing

When the feature is a **numeric variable**, we can conduct missing data imputation.

We replace the missing values with the **average** or **median** value from the data of the same feature that is not missing.

When the feature is a **categorical variable**, we may impute the missing data by the **mode** (the most frequent value)

	col1	col2	col3	col4	col5		col1	col2	col3	col4	col5	
0	2	5.0	3.0	6	NaN	mean()	0	2.0	5.0	3.0	6.0	7.0
1	9	NaN	9.0	0	7.0		1	9.0	11.0	9.0	0.0	7.0
2	19	17.0	NaN	9	NaN		2	19.0	17.0	6.0	9.0	7.0

### Solution : Replace the Missing

For **categorical features**, we can add a new category with a value such as “**\_MISSING\_**”.

For **numerical features**, we can replace it with a particular value such as **-999**.

This way, we are still keeping the missing values as valuable information.



shutterstock.com • 1049623214

### Irregular data (Outliers)

**Outliers** are data that is **distinctively different** from other observations.

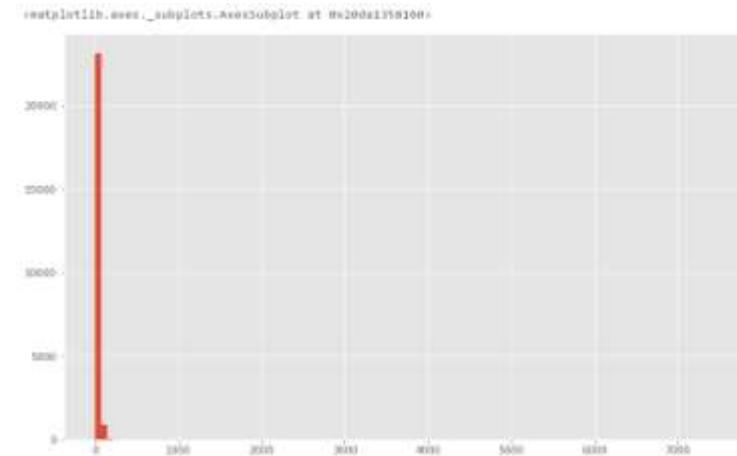
They could be **real outliers** or **mistakes**.



### Technique : Histogram/Box Plot

When the feature is numeric, we can use a histogram and box plot to detect outliers.


The **data looks highly skewed** with the possible existence of outliers.





To study the feature closer, let's make a **box plot**.

In this plot, we can see there is an outlier at a **value of over 7000**.



### Technique : Descriptive Statistics

For **numeric features**, the outliers could be too distinct that the **box plot can't visualize** them.

Instead, we can look at their **descriptive statistics**.

For example, for the feature *life\_sq* again, we can see that the **maximum value is 7478**, while the **75% quartile is only 43**.

The 7478 value is an **outlier**.

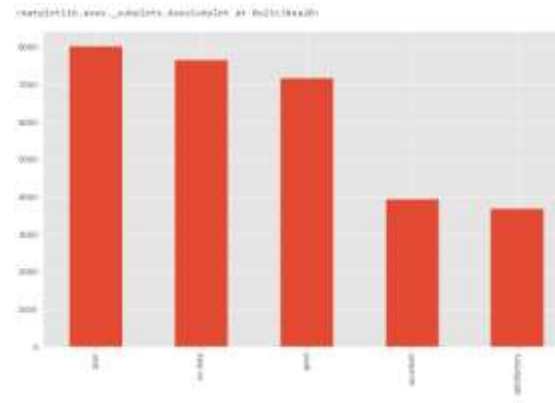
```
count    24088.000000
mean      34.403271
std       52.285733
min        0.000000
25%       20.000000
50%       30.000000
75%       43.000000
max       7478.000000
Name: life_sq, dtype: float64
```

### Technique : Bar Chart

When the feature is **categorical**, we can use a bar chart to learn about its categories and distribution.

For example, the feature *ecology* has a reasonable distribution.

But if there is a category with **only one value** called “other”, then that would be an **outlier**.



After all the hard work done for missing data and outliers, let's look at **unnecessary data**, which is more straightforward.

The unnecessary data is when the data **doesn't add value**.

We cover three main types of unnecessary data due to different reasons.

### Unnecessary type : Uninformative / Repetitive

Sometimes one feature is uninformative because it has too many rows being the same value.

### How to find out?

We can create a list of features with a high percentage of the same value.

For example, we specify below to show features with over 95% rows being the same value.

```
oil_chemistry_raion: 99.02858%
no      34175
yes      296
Name: oil_chemistry_raion, dtype: int64

railroad_terminal_raion: 96.22187%
no      29335
yes      1136
Name: railroad_terminal_raion, dtype: int64

nuclear_reactor_raion: 97.16788%
no      29608
yes      863
Name: nuclear_reactor_raion, dtype: int64

big_road1_iline: 97.43891%
no      29699
yes      781
Name: big_road1_iline, dtype: int64

railroad_iline: 97.06934%
no      29578
yes      893
Name: railroad_iline, dtype: int64

cafe_count_500_price_high: 97.26641%
0      29635
1       787
2        88
3         11
Name: cafe_count_500_price_high, dtype: int64

mosque_count_500: 99.51101%
0      30322
1       149
Name: mosque_count_500, dtype: int64

cafe_count_1000_price_high: 95.52889%
0      29184
1       1184
2        145
3         51
4          9
5          5
6          8
7          1
Name: cafe_count_1000_price_high, dtype: int64

mosque_count_1000: 98.08342%
0      29882
1        584
Name: mosque_count_1000, dtype: int64

mosque_count_1500: 96.21936%
0      29319
1       1152
Name: mosque_count_1500, dtype: int64
```

### Unnecessary Type :

- Irrelevant

- Duplicates

```
timestamp  full_sq  life_sq  floor  build_year  num_room  price_doc
2014-12-09  40      -999.0  17.0  -999.0      1.0      4607265    2
2014-04-15  134      134.0  1.0    0.0        3.0      5798496    2
2013-08-30  40      -999.0  12.0  -999.0      1.0      4462000    2
2012-09-05  43      -999.0  21.0  -999.0     -999.0    6229540    2
2013-12-05  40      -999.0  5.0    -999.0      1.0      4414080    2
2014-12-17  62      -999.0  9.0    -999.0      2.0      6552000    2
2013-05-22  68      -999.0  2.0    -999.0     -999.0    5406690    2
2012-08-27  59      -999.0  6.0    -999.0     -999.0    4506800    2
2013-04-03  42      -999.0  2.0    -999.0     -999.0    3444000    2
2015-03-14  62      -999.0  2.0    -999.0      2.0      6520500    2
2014-01-22  46      28.0   1.0    1968.0     2.0      3000000    2
2012-10-22  61      -999.0  18.0  -999.0     -999.0    8248500    2
2013-09-23  85      -999.0  14.0  -999.0      3.0      7725974    2
2013-06-24  40      -999.0  12.0  -999.0     -999.0    4112800    2
2015-03-30  41      41.0   11.0   2016.0     1.0      4114580    2
2013-12-18  39      -999.0  6.0    -999.0      1.0      3700946    2
2013-08-29  58      58.0   13.0   2013.0     2.0      5764128    1
          50      33.0   2.0    1972.0     2.0      8150000    1
          52      30.0   9.0    2006.0     2.0     10000000    1
2013-08-30  38      17.0   15.0   2004.0     1.0      6400000    1
Name: id, dtype: int64
```

There are 16 duplicates based on this set of key features.

### Inconsistent : Capitalization

#### What to do?

To avoid this, we can put all letters to lower cases (or upper cases).

```
Poselenie Sosenskoe      1776
Nekrasovka               1611
Poselenie Vnukovskoe     1372
Poselenie Moskovskij      925
Poselenie Voskresenskoe   713
...
Molzhaninovskoe          3
Poselenie Kievskij        2
Poselenie Shhapovskoe     2
Poselenie Mihajlovo-Jarcevscoe 1
Poselenie Klenovskoe       1
Name: sub_area, Length: 146, dtype: int64
```

# STATISTICS FOR DATA SCIENCE

## Inconsistent Data



### Inconsistent Type : Formats

Another standardization we need to perform is the data formats.

One example is to convert the feature from string to DateTime format.

### How to find out?

The feature *timestamp* is in string format while it represents dates.

	id	timestamp	lat	lon	floor	mag	mag_floor	material	bed_jen	non_poor	bed_jen	lat	lon	sub_point_200	price_high	big_church_point_200	church_point_200	mosque_point_200	house_point_200	spot_point_200	market_point_200	price_low	sub_area_low	category
0	1	2014-02-01	42	21	40	160	160	160	160	160	160	160	160	0	15	20	1	0	0	14	50000	Wairoa	godun_jen	
1	2	2014-02-01	34	10	30	160	160	160	160	160	160	160	160	0	15	20	1	0	0	14	10000	Ngahia	jen	godun_jen
2	3	2014-02-01	40	10	30	160	160	160	160	160	160	160	160	0	10	20	0	0	4	0	10	50000	eastern	jen
3	4	2014-02-01	40	10	30	160	160	160	160	160	160	160	160	1	4	4	0	0	0	20	1	10000	Wairoa	godun_jen
4	5	2014-02-01	7	10	40	160	160	160	160	160	160	160	160	11	0	0	20	0	0	0	14	10000	Wairoa	godun_jen
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
3466	3466	2014-02-01	44	21	10	160	160	160	160	160	160	160	160	0	15	20	1	0	0	14	0	10000	Wairoa	godun_jen
3467	3470	2014-02-01	40	10	30	160	160	160	160	160	160	160	160	24	0	0	10	1	0	0	15	10000	Wairoa	jen
3468	3471	2014-02-01	40	10	30	160	160	160	160	160	160	160	160	0	1	0	0	0	1	11	1	10000	Wairoa	jen
3469	3472	2014-02-01	44	21	10	160	160	160	160	160	160	160	160	1	4	0	1	0	0	14	1	10000	Wairoa	jen
3470	3473	2014-02-01	40	10	30	160	160	160	160	160	160	160	160	0	1	0	0	0	0	14	10	10000	Wairoa	jen

347 rows, 24 columns



### What to do?

We can convert it and extract the date or time values by using the code below. After this, it's easier to analyze the transaction volume group by either year or month.

```
2014    13662
2013     7978
2012     4839
2015     3239
2011      753
Name: year, dtype: int64
```

```
12    3408
4     3191
3     2972
11    2978
10    2736
6     2578
5     2496
9     2346
2     2275
7     1875
8     1831
1     1809
Name: month, dtype: int64
```


### Inconsistent Type : Categorical Values

Inconsistent categorical values are the last inconsistent type we cover.

A categorical feature has a limited number of values. Sometimes there may be other values due to reasons such as typos.

### How to find out?

For instance, the value of *city* was typed by mistakes as “torontoo” and “tronto”. But they both refer to the correct value “toronto”.



### Inconsistent Type : Addresses

The address feature could be a headache for many of us. Because people entering the data into the database often *don't* follow a standard format.

### How to find out?

We can find messy address data by looking at it. Even though sometimes we can't spot any issues, we can still run code to standardize them.

	address
0	123 MAIN St Apartment 15
1	123 Main Street Apt 12
2	543 FirSt Av
3	876 FIRst Ave.

	address	address_std
0	123 MAIN St Apartment 15	123 main st apt 15
1	123 Main Street Apt 12	123 main st apt 12
2	543 FirSt Av	543 first ave
3	876 FIRst Ave.	876 first ave



# THANK YOU

---

**D. Uma**

Department of Computer Science and Engineering

**[umaprabha@pes.edu](mailto:umaprabha@pes.edu)**