



BIG DATA

MapReduce Algorithms

– Matrix Multiplication

K V Subramaniam

Computer Science and Engineering

Matrices and Vectors - introduction



Matrix Multiplication algorithms

Fundamental to many computations, including
Page Rank

Source

Leskovec, Jure, Anand Rajaraman, and Jeffrey
David Ullman. *Mining of massive datasets*.
Cambridge University Press, 2014.

[http://infolab.stanford.edu/~ullman/mmds/book.
pdf](http://infolab.stanford.edu/~ullman/mmds/book.pdf)

4.3.2 of T1

Vectors

Can be defined as an ordered list of numbers

Visualization

An arrow where the direction of the vector is given by the relative size of the components


Some Common Operations

Addition: $\mathbf{v} + \mathbf{w}$

Add components

Scalar multiplication $a\mathbf{v}$

Multiply each component by constant



Rectangular array of numbers.

The numbers are called the elements of the matrix.

An $m \times n$ matrix has m rows and n columns

Can be considered as a collection of
 m row vectors


n column vectors

An $n \times n$ matrix is called a square matrix.

Vector can be considered as a

$1 \times n$ matrix (row matrix)

$n \times 1$ matrix (column matrix)



Each element of a matrix is often denoted by a variable with two **subscripts**. For example, $a_{2,1}$ represents the element at the second row and first column of a matrix **A**.

Multiply each row vector of **A** by the corresponding elements of **x** and sum

Multiplying a $m \times n$ matrix by an n element vector gives an m element vector

$$A\mathbf{x} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$
$$= \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \cdots + a_{1n}x_n \\ a_{21}x_1 + a_{22}x_2 + \cdots + a_{2n}x_n \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \cdots + a_{mn}x_n \end{bmatrix}.$$

$$A = \begin{bmatrix} 1 & 0 & -1 \\ 3 & 1 & 2 \end{bmatrix}, \quad \mathbf{x} = (x, y, z)$$

$$A\mathbf{x} = \begin{bmatrix} 1 & 0 & -1 \\ 3 & 1 & 2 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} x - z \\ 3x + y + 2z \end{bmatrix}$$
$$= (x - z, 3x + y + 2z).$$

Typically, matrices are stored as multi-dimensional arrays in programs

```
int A[10][10]
```

allocates 100 integers and is accessed as a 10x10 matrix

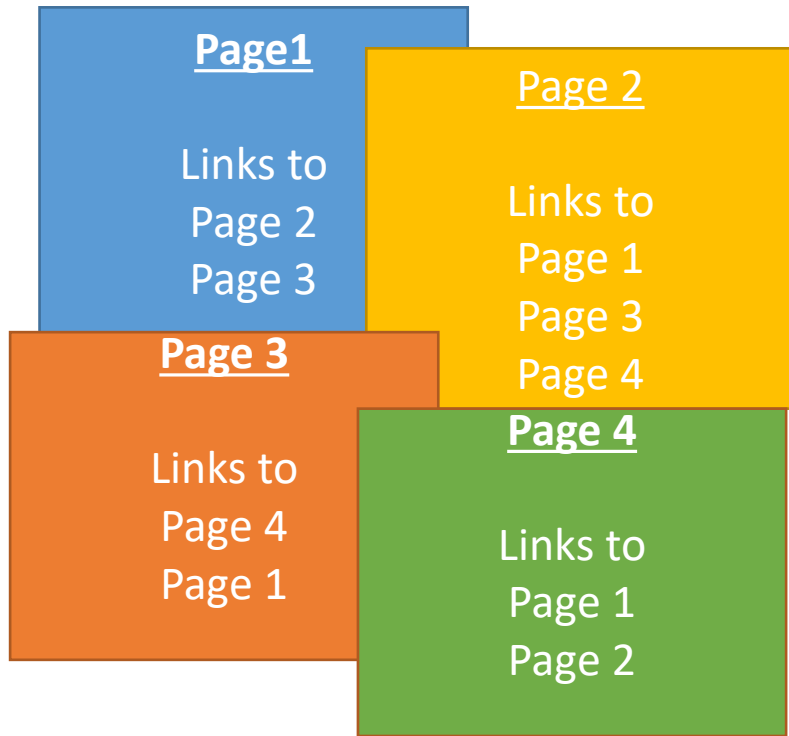
Space required to store the matrix – =
 $10 \times 10 \times \text{sizeof}(\text{int}) = 100 \times \text{sizeof}(\text{int}) = 100 \times 4 = 400$ bytes.

In general, we need n^2 integers to store an ***nxn*** matrix

Matrix Representation of WWW

BIG DATA

How do you represent the pages in WWW?


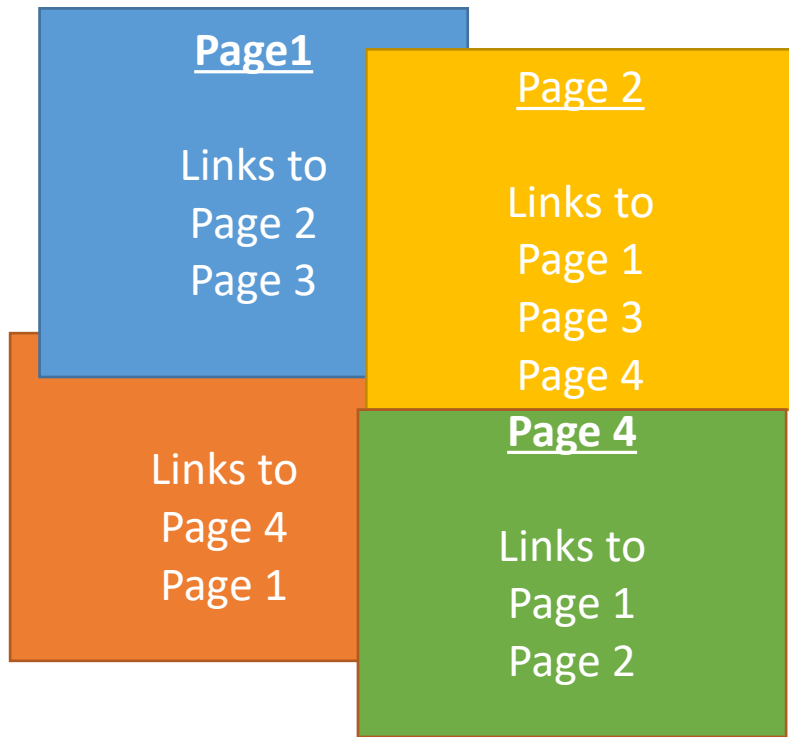


Consider a sample of the internet that contains 4 pages

- How should we represent this?

BIG DATA


Modelling the WWW as a directed graph




Represented as a directed graph

Pages in the WWW

Representing the graph as an Adjacency Matrix



Directed Graph




This is fine for a small graph – 4 pages.

But internet is large – billions of pages.


How much storage will we require?

Large scale matrix representations




BIG DATA

Representing the graph as an Adjacency Matrix



Directed Graph



Internet is large – billions of pages.

Note – most of the entries will be 0

Store as a sparse matrix..

BIG DATA

Sparse Matrix representation



In Big Data, we deal with large matrices
e.g, n will be the order of 10^{10} if n
is number of web pages

And it will be a sparse matrix

W o n t f i t i n t h e m e m o r y D R A M

Have to store it in HDFS

BIG DATA

HDFS Sparse matrix representation



Store only non-zero elements as a separate record in CSV format

For each element store
<row_number, column_number, value>
As the format

As many entries as there are links

Exercise –Store the graph given on the right into a HDFS CSV file

0	1	1	1
1	0	0	1
1	1	0	0
0	1	1	0

BIG DATA

Solution

1, 2, 1
1, 3, 1
1, 4, 1
2, 1, 1
2, 4, 1
3, 1, 1
3, 2, 1
4, 2, 1
4, 3, 1

0	1	1	1
1	0	0	1
1	1	0	0
0	1	1	0

As an exercise, try saving this in a file and loading it onto HDFS that you have installed.

Matrix Vector Multiplication

To multiply an $n \times n$ matrix M with an n -element vector v , compute

$$x_i = \sum_{j=1}^n m_{ij} v_j$$

$$A\mathbf{x} = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n \\ a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \dots + a_{mn}x_n \end{bmatrix}$$

Let us assume that the vector \mathbf{v} fits into memory

Vector \mathbf{v} is shared by all the mappers

M_{ij} is stored as a CSV file on HDFS and is distributed across multiple nodes

$$x_i = \sum_{j=1}^n m_{ij} v_j$$

$$A\mathbf{x} = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n \\ a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \dots + a_{mn}x_n \end{bmatrix}$$

← 2

$$x_i = \sum_{j=1}^n m_{ij} v_j$$

map:


Computes the partial product

Uses the key as i the index into the target vector

output $(i, m_{ij} v_j)$

reduce:

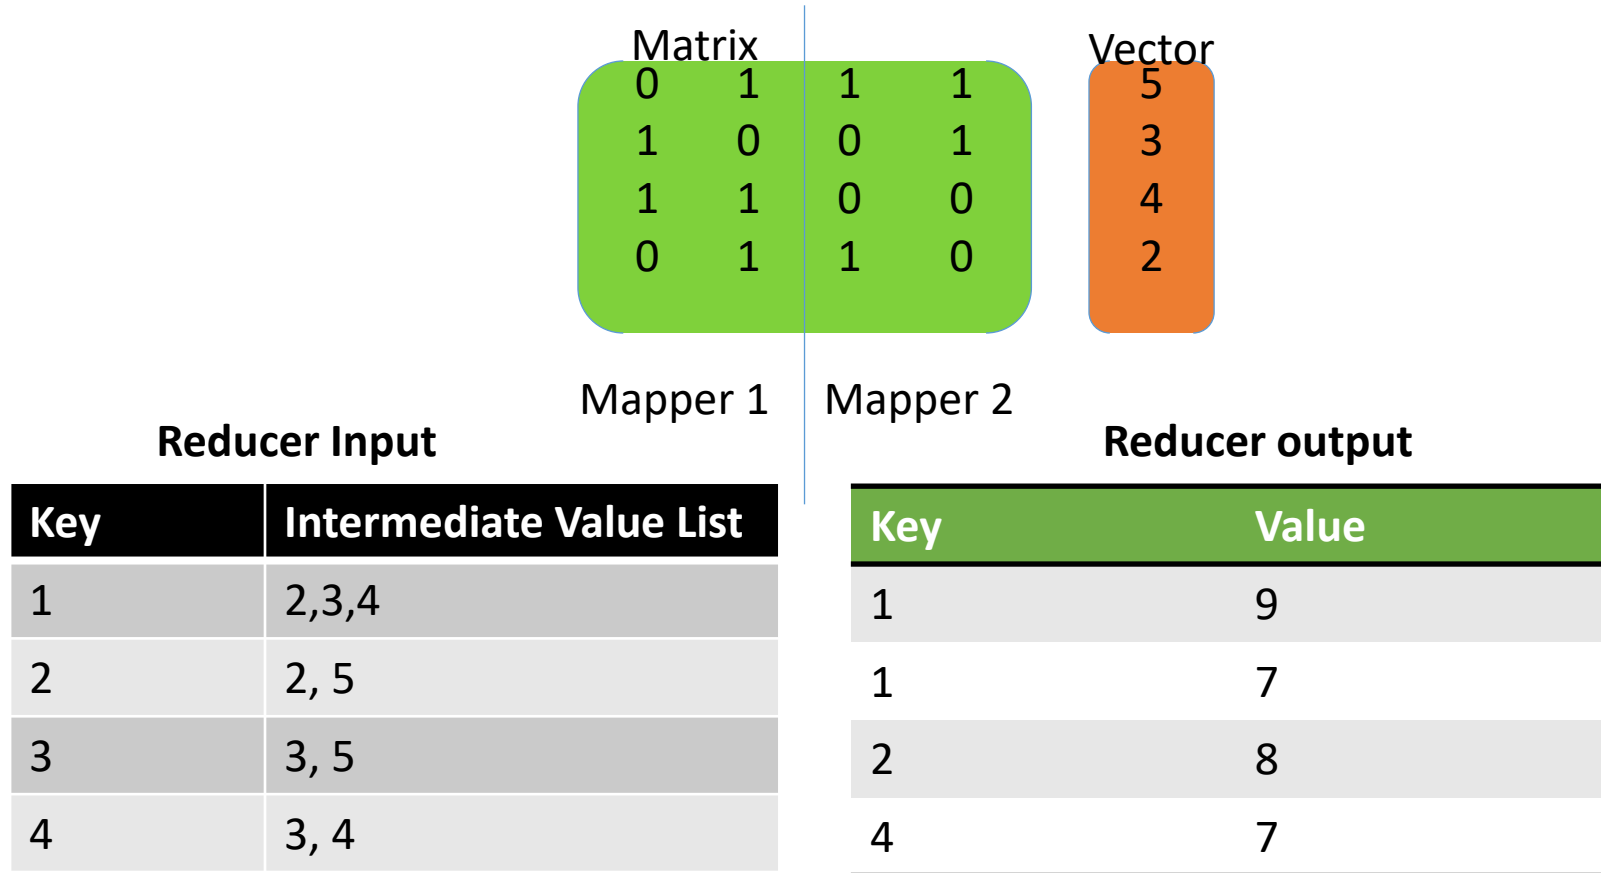
Sums all the partial products.



Note that the key is the index into the vector where this value will contribute

Key	Value
1	$1 * 3 = 3$
2	$1 * 5 = 5$
3	$1 * 5 = 5$
3	$1 * 3 = 3$
4	$1 * 3 = 3$

Key	Value
1	$1 * 4 = 4$
1	$1 * 2 = 2$
2	$1 * 2 = 2$
4	$1 * 4 = 4$



Matrix

0	0	3	8
0	9	5	0
0	10	0	0
5	0	0	1

Vector

1
2
3
4

Assuming rows 1 and 3 are in datanode1 and 2 and 4 are in datanode 2, perform a matrix multiplication using Map Reduce

Show inputs/outputs of mappers/reducers as discussed in previous slides

Matrix Vector Multiplication - extensions

Case 2: v does not fit into main memory.

Partition M and v into stripes.




Figure 2.4: Division of a matrix and vector into five stripes

The same MapReduce algorithm can be used.



THANK YOU

K V Subramaniam

Dept. of Computer Science and Engineering

subramaniamkv@pes.edu,

ushadevibg@pes.edu