



**BIG DATA**

**Big Data Algorithms**

**– Relational Operations**

---

**K V Subramaniam**

Computer Science and Engineering

# BIG DATA

## Class Overview

---

- Relational algebra overview
- Select and Project with MR
- Set Operations
- Join
- Grouping and Aggregation
- Case Study: HIVE



## Relational Operations

- Relations
  - Tables; columns = attributes
  - Rows = tuples;  $R(A_1, A_2, \dots, A_n)$
- Relational Operators
  - Selection:  $\sigma_C(R)$  select from  $R$  according to condition  $C$
  - Projection  $\pi_S(R)$  select from  $R$  subset of attributes  $S$
  - Union, intersection, difference
  - Natural join
  - Grouping: partition  $R$  according to attributes  $G$
  - Aggregation:  $SUM, COUNT, AVG, MAX, MIN$

# BIG DATA

## Simple Problem



id	Name	Role	Team
1	Virat Kohli	Captain	RCB
2	Gautham Gambhir	Captain	KKR
3	Anil Kumble	Coach	MI
4	Virender Sehwag	Coach	KXIP

1. Write an SQL query to list
  - a) All the details for the Coaches.
  - b) Only the names of the coaches
  - c) Total #coaches.
2. What type of a relational operation are we using?

# BIG DATA

## Simple Problem

---



- Data in a database is well structured.
- Instead assume that data is stored in a file in HDFS as shown below.

```
1, Virat Kohli, Captain, RCB  
2, Gautham Gambhir, Captain, KKR  
3, Anil Kumble, Coach, MI  
4, Virender Sehwag, Coach, KXIP
```

- Now we need to perform the query.

## Select and Project in MapReduce

- map
  - Read each row  $t$  of table
  - Check if it satisfies condition  $C$
  - If so, output  $(t,t)$
- Reduce
  - do nothing

How will the map output look for the query 1a ?

Map

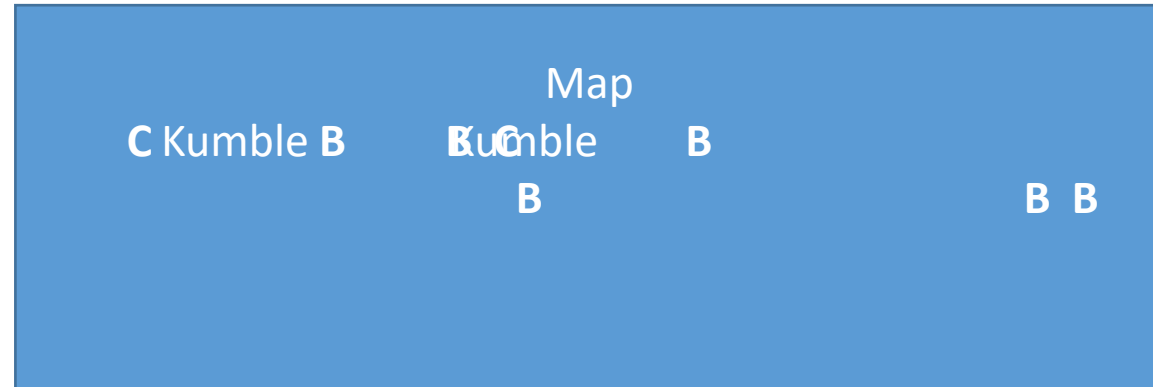
3, Anil Kumble, Coach, MI	B	B	B	B
4, Virender Sehwag, Coach, KXIP	B	B	B	B

Reduce

3, Anil Kumble, Coach, MI  
4, Virender Sehwag, Coach, KXIP



- map
  - Read each row  $t$  of table
  - Calculate subset of attributes  $t$
  - Output  $(t, t)$
- reduce
  - Eliminate duplicates
    - $(t, t) \rightarrow (t, t)$






## **Relational Operations requiring two input files – Set operations**



# BIG DATA

## Class Exercise - Union...

- Union  $R \cup S$ 
  - $R$  and  $S$  have the same structure
  - Output records in  $R$  or  $S$
  - *We need to solve a problem here*
- The basic problem :
  - Need to read in 2 input files
  - produce 1 output file
- Reading multiple input files
  - Same problem in Matrix-vector multiplication
  - Need to read in matrix  $M$  and vector  $v$



# BIG DATA

## One Solution

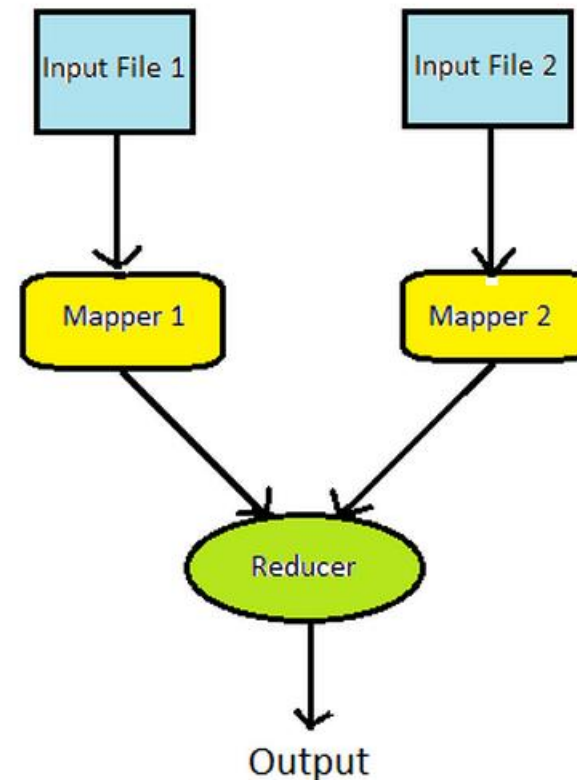
---

- map reads in one file
- If we have multiple input files, we can try to combine them into one file
  - This can be used as input to map
- Problem: multiple input files can be combined in many ways
  - We can have all the records of one file followed by all the records of another file
  - Or merge the two files and sort the records
  - Files may not have the same structure


# BIG DATA

## The Hadoop Solution


- Use map reduce itself to merge the files
- Example
  - We have two mappers
  - Each mapper reads in one file
  - Each mapper writes a key-value pair
  - The reducer uses the keys to merge the files
  - This output can be used as input to subsequent map reduce passes



- Union *R U S*
  - Need to read in 2 input files and produce 1 output file
  - Need to have 2 mappers reading different input files
  - Can be done using *MultipleInput* option in Hadoop
    - MultipleInputs.
      - addInputPath (job,
      - new Path (args[0]),
      - TextInputFormat.class,
      - Mapper1.class);



- mapper 1
  - Read each row of table  $R$
  - Output  $(t,t)$
- mapper 2
  - Read each row of table  $S$
  - Output  $(t,t)$
- reducer
  - Eliminate duplicates if any
    - $(t,[t,t]) \rightarrow (t,t)$



## Sample Problem

---


- Given the following input for two files
  - File 1
    - A
    - B
    - C
    - D
  - File 2
    - A
    - E
    - F
    - C

- Show (for the Union algorithm)
  - Input and output of mapper 1
  - Input and output of mapper 2
  - Input and output of reducer




- Mapper 1 output
    - A , A
    - B, B
    - C, C
    - D, D
  - Mapper 2 output
    - A, A
    - E, E
    - F, F
    - C, C
- Reducer Input
    - A, [A, A]
    - B, B
    - C, [C, C]
    - D, D
    - E, E
    - F, F
  - Reducer Output
    - A
    - B
    - C
    - D
    - E
    - F

- Compute  $R \cap S$
- mapper 1
  - Read each row of table  $R$
  - Output  $(t, t)$
- mapper 2
  - Read each row of table  $S$
  - Output  $(t, t)$
- reducer
  - Output only duplicates
    - $(t, [t, t]) \rightarrow (t, t)$



- Compute  $R - S$ 
  - All rows in  $R$ , not in  $S$
- mapper 1
  - Read each row of table  $R$
  - Output  $(t, R)$
- mapper 2
  - Read each row of table  $S$
  - Output  $(t, S)$
- reducer
  - Output only
    - $(t, [R]) \rightarrow (t, t)$




**Join**

# BIG DATA

## Natural Join

- Join R and S on attributes B
  - A, C are the other attributes in R,S
- mapper 1
  - Read (a,b) of R, output (b,(R,a))
- mapper 2
  - Read (b,c) of S, output (b,(S,c))
- reducer
  - for each pair (b,(R,a)) and (b,(S,c)), output (a,b.c)



- Given the following input for two files
  - Table Employee E(Name, age)
    - Gabbar 35
    - Viru 37
    - Jai 33
    - Baldev 44
    - Basanti 31
  - Table Dept D(Name, Dept)
    - Gabbar Bandit
    - Viru Hero
    - Jai Hero
    - Baldev Police
    - Basanti Heroine

- Show (for the Natural Join algorithm)
  - Input and output of mapper 1
  - Input and output of mapper 2
  - Input and output of reducer

- Mapper 1 Output
  - Gabbar, (E, 35)
  - Viru, (E, 37)
  - Jai, (E, 33)
  - Baldev, (E, 44)
  - Basanti, (E, 31)
- Mapper 2 Output
  - Gabbar, (D, Bandit)
  - Viru, (D, Hero)
  - Jai, (D, Hero)
  - Baldev (D, Police)
  - Basanti (D, Heroine)
- Reducer Input
  - Gabbar, (E, 35), (D, Bandit)
  - Viru, (E, 37), (D, Hero)
  - Jai, (E, 33), (D, Hero)
  - Baldev, (E, 44), (D, Police)
  - Basanti, (E, 31), (D, Heroine)
- Reducer Output
  - Gabbar, 35, Bandit
  - Viru, 37, Hero
  - Jai, 33, Hero
  - Baldev, 44, Police
  - Basanti, 31, Heroine)

•

## Grouping and Aggregation



# BIG DATA

## Grouping and Aggregation...

---

- For relation  $R(A,B,C)$  group by A and aggregate by function  $f(B)$ 
  - Social networking site has a relationship *Friends* (*User*, *Friend*, *Date of friendship*)
  - Grouping by *User* and aggregating by *COUNT* (*Friend*) produces a table
    - First column is the *User*
    - Second column is the count of *Friends* for the *User*

# BIG DATA

## ...Grouping and Aggregation

---

- map
  - for each line (a,b,c) (e.g., User, Friend, Date)
  - Output (a,b) (e.g., User, Friend)
- reduce
  - Aggregate (a, [b1 , b2 , b3 , ...]) into (a, f(b1 , b2 , b3 , ...))

## Case Study : HIVE

# BIG DATA

## What is HIVE

---

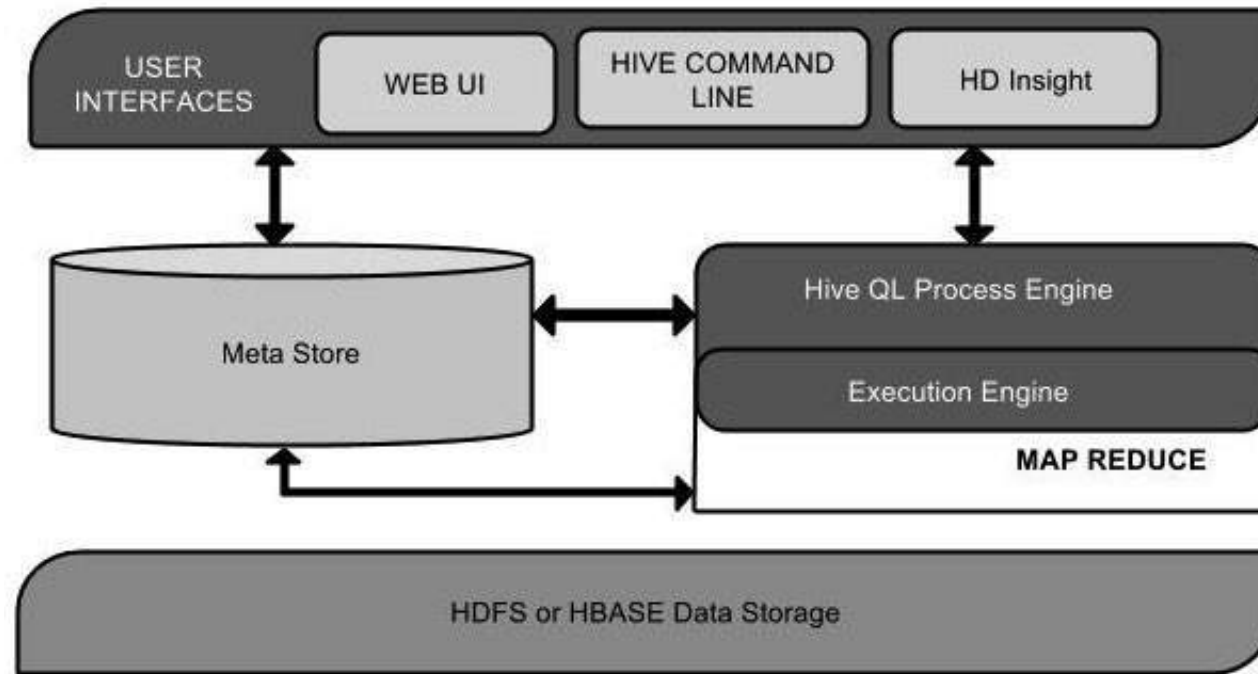


- A system for querying and managing structured data built on top of Map/Reduce and Hadoop
- Facebook data
  - Structured logs with rich data types (structs, lists and maps)
  - A user base wanting to access this data in the language of their choice
  - A lot of traditional SQL workloads on this data (filters, joins and aggregations)
  - Other non SQL workloads

# BIG DATA

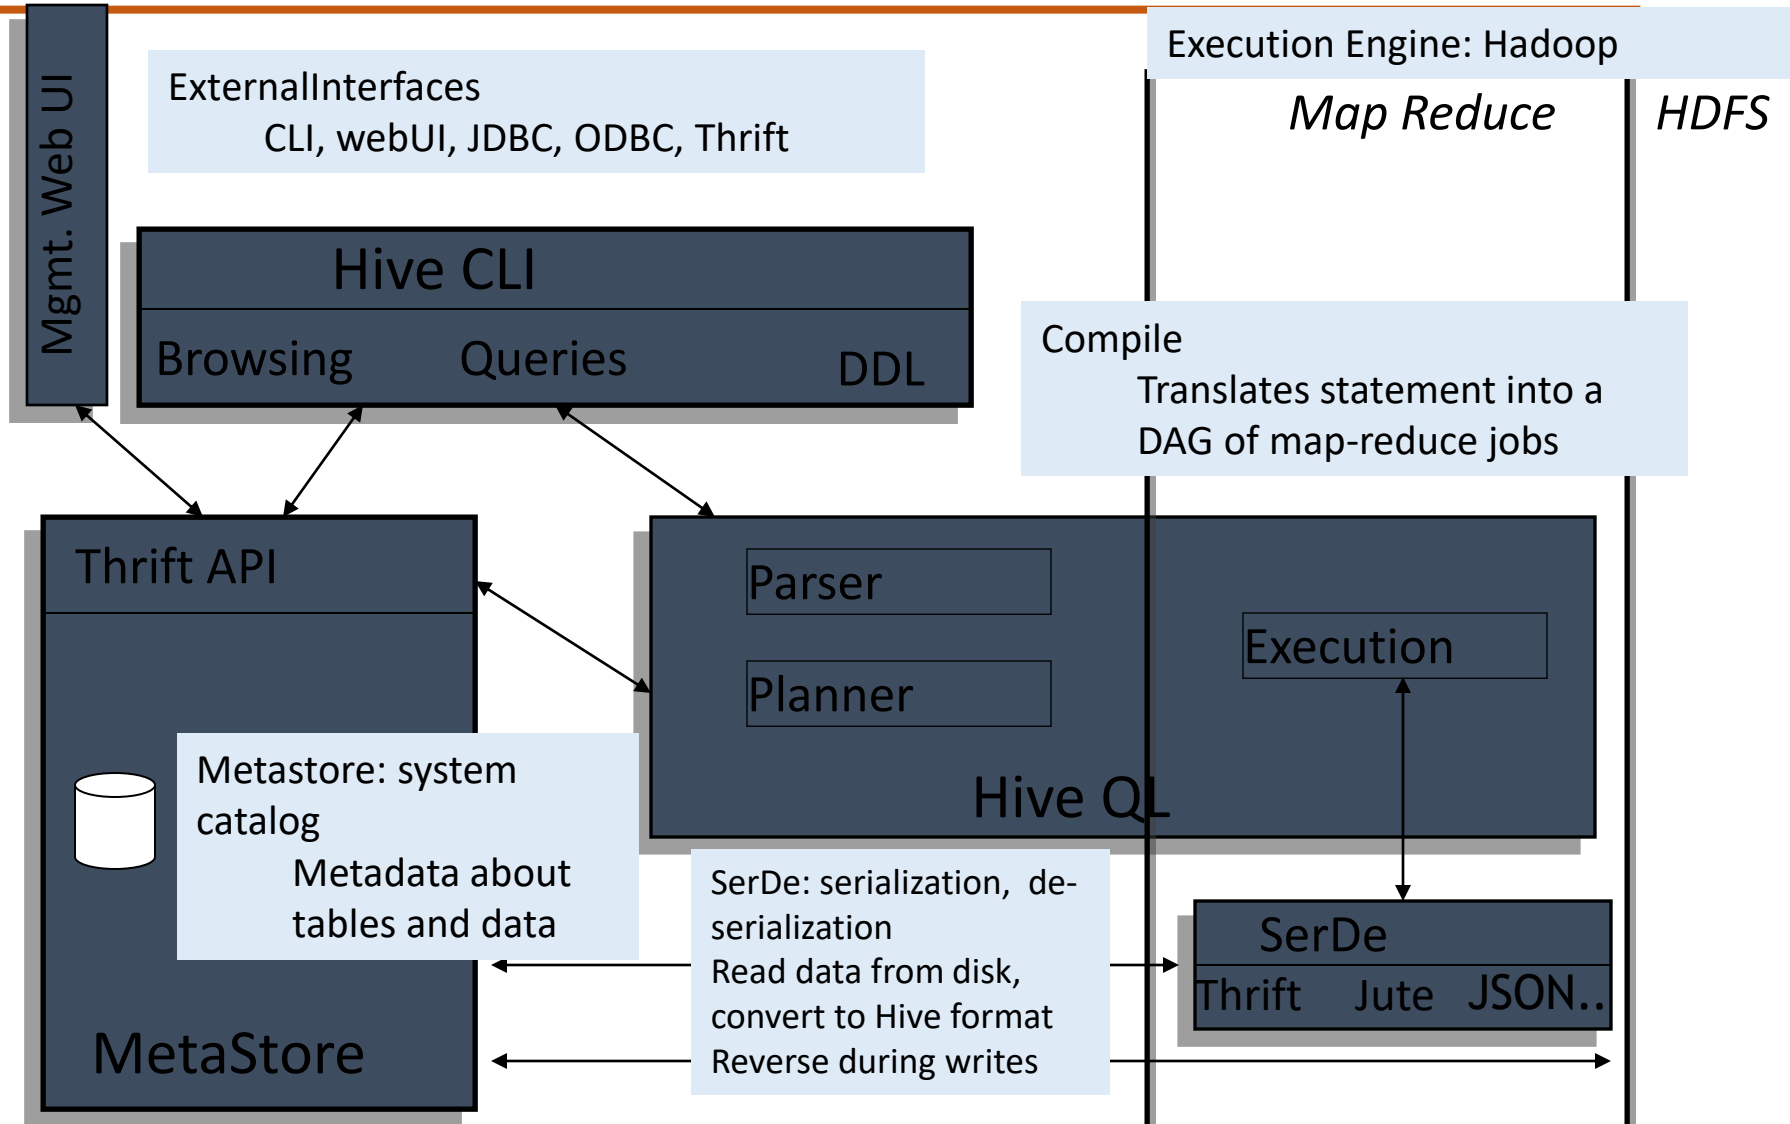
## What is HIVE

- Data is stored in HDFS
- Meta Store for schema
- User submits SQL query
  - Converted to MR jobs



# BIG DATA

## HIVE Components



### 1. Hive data stored as HDFS files

/hive/clicks

/hive/clicks/ds=2008-03-25

/hive/clicks/ds=2008-03-25/0

2. Table similar to table in relational database  
Mapped to HDFS directory  
Data for table clicks is in the directory  
/hive/clicks  
For scalability, tables divided into multiple  
files and directories

### 3. Partition

Part of table partitioned on values of columns

Implemented as a n HDFS subdirectory

If clicks is partitioned on column ds Data with a  
particular ds value 2008-03-25 will be stored in  
/hive/clicks/ds=2008-03-25

If further divided on ctry value US will be stored in  
the directory /hive/clicks/ds=2008-03-25/ctry=US.

### 4. Bucket

Subdivision of partition

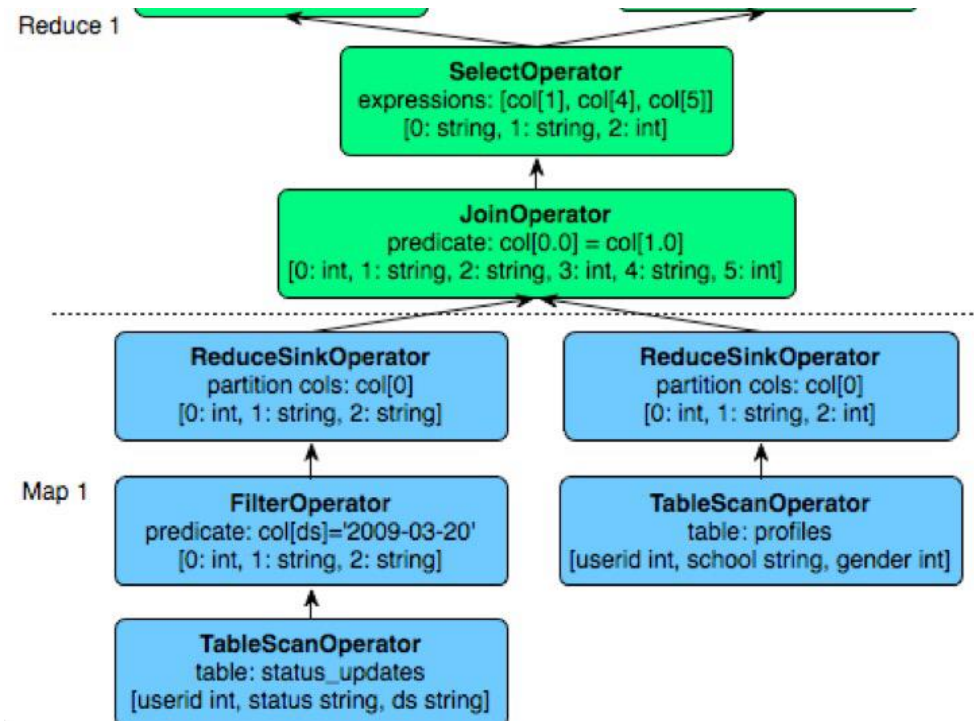
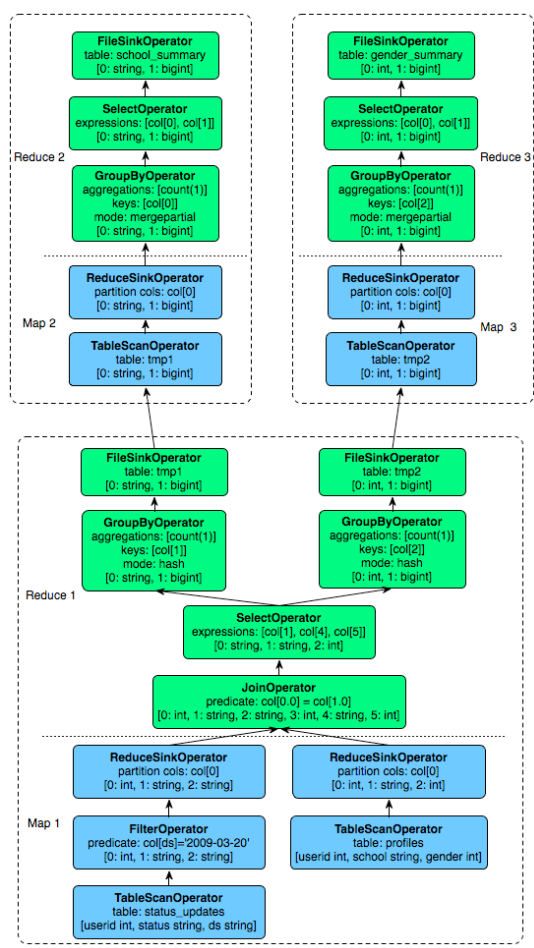
Divided based on hash of a specified column

Thusoo, Ashish, Joydeep Sen Sarma, Namit Jain, Zheng Shao, Prasad Chakka, Suresh Anthony, Hao Liu, Pete Wyckoff, and Raghotham Murthy. "Hive: a warehousing solution over a map-reduce framework." *Proceedings of the VLDB Endowment* 2, no. 2 (2009): 1626-1629

```
FROM (SELECT a.status, b.school, b.gender
      FROM status_updates a JOIN profiles b
      ON (a.userid = b.userid and
          a.ds='2009-03-20' )
      ) subq1
INSERT OVERWRITE TABLE gender_summary
      PARTITION(ds='2009-03-20')
SELECT subq1.gender, COUNT(1) GROUP BY subq1.gender
INSERT OVERWRITE TABLE school_summary
      PARTITION(ds='2009-03-20')
SELECT subq1.school, COUNT(1) GROUP BY subq1.school
```



```
FROM (SELECT a.status, b.school, b.gender
      FROM status_updates a JOIN profiles b
      ON (a.userid = b.userid and
          a.ds='2009-03-20' )
    ) subq1
```





# THANK YOU

---

**K V Subramaniam**

Dept. of Computer Science and Engineering

**[subramaniamkv@pes.edu](mailto:subramaniamkv@pes.edu)**